

# LOGIC FILE SYSTEM

the filesystem of the 21st century

or hierarchical filesystems considered harmful

YOANN PADIOLEAU

supervised by OLIVIER RIDOUX

IRISA / University of Rennes 1

# Plan

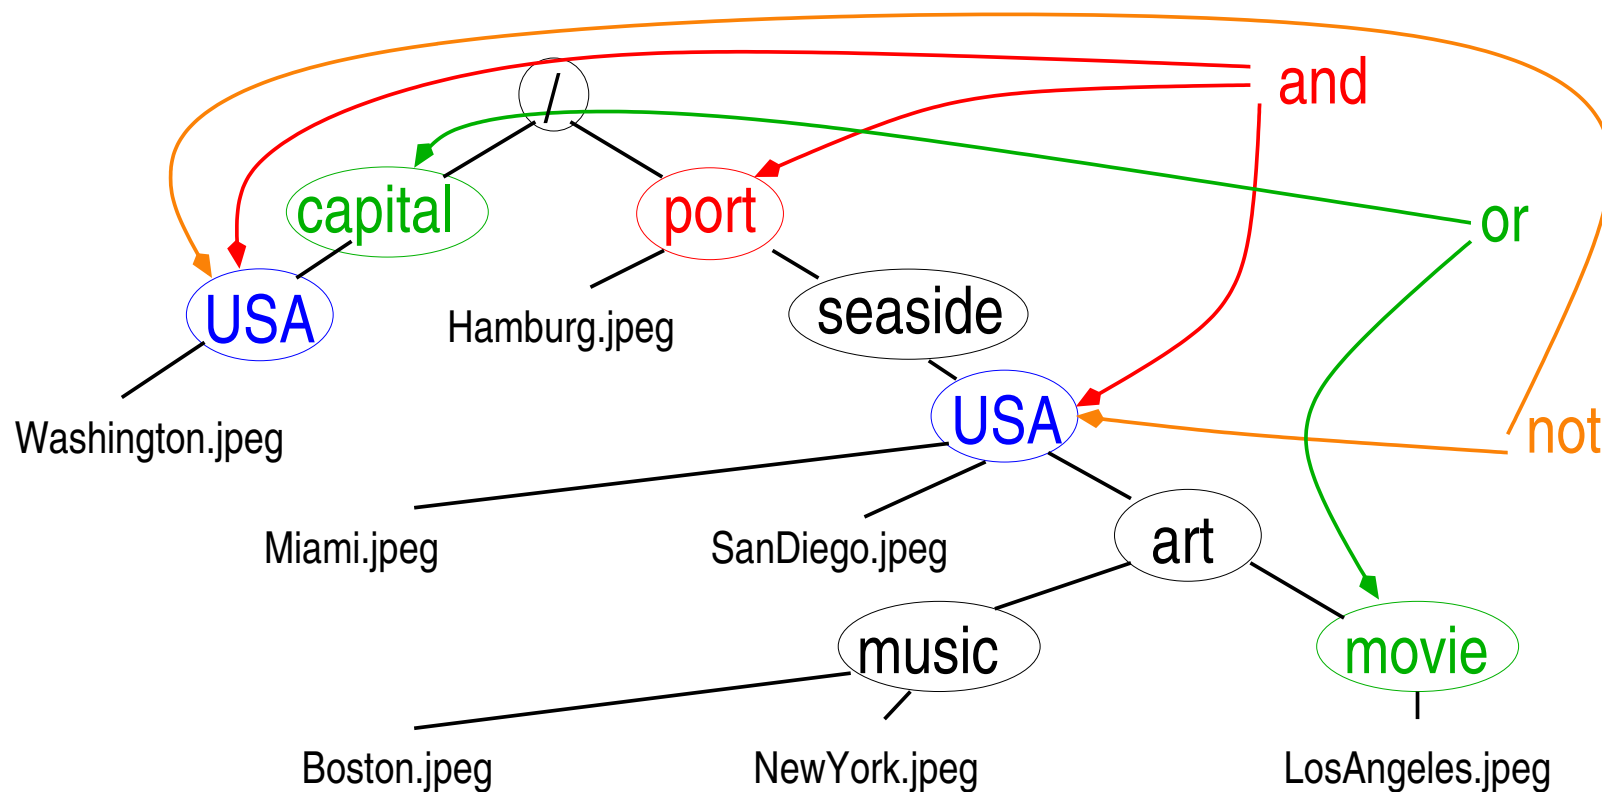
- Motivations  
(combine navigation and query in a filesystem: advanced dirs)  
(several points of view on file contents: advanced files)
- Demo (managing photos, music, code, bibtex, ...)
- Specification (cd = ?, ls = ?, mv = ?, cat = ?, ...)
- Security (read-write-execute = ?, user-group-other = ?)
- Implementation and evaluation (fast enough ?)
- Related work (advanced filesystems, search tools, and IDE)
- Conclusion and further works (take over the world)

# A toy example

to represent a collection of city maps

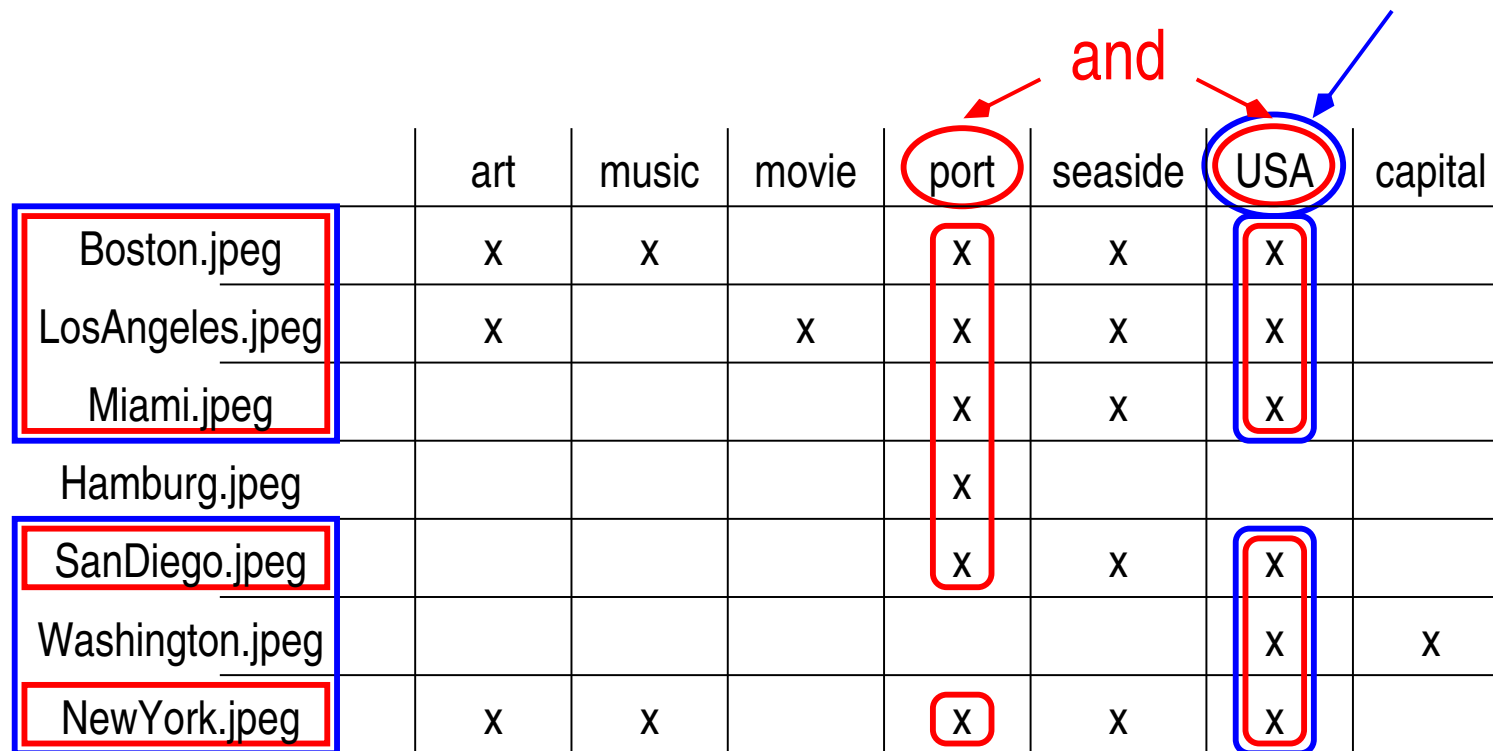
- a collection of cities — Boston, Hamburg, Los Angeles, Miami, New York, San Diego, Washington
- a collection of descriptive attributes
  - to be a port, on the seaside,
  - to be in the USA,
  - to be a capital,
  - to be an art city, for music, or movie

# A hierarchical filesystem



a meaning for **cd /USA**, **cd not USA**, **cd capital or movie**, and **cd port and USA**?

# Boolean organization (Google)



	art	music	movie	port	seaside	USA	capital
Boston.jpeg	x	x		x	x	x	
LosAngeles.jpeg	x		x	x	x	x	
Miami.jpeg				x	x	x	
Hamburg.jpeg				x			
SanDiego.jpeg				x	x	x	
Washington.jpeg						x	x
NewYork.jpeg	x	x		x	x	x	

good for cd not USA and cd capital or movie

but, not incremental enough for cd USA and cd port and USA

# Motivation

→ merge navigation  
and querying  
in a file system

(as in hierarchical organizations)




(as in boolean organizations)




(every tool benefits of it: from shells to multimedia players)




advanced (virtual) directories

# LFS organization

- assign multiple properties to files


file1.txt =   

file2.jpg =   


file7.jpg =   

⋮

an example:  
Venn diagram  
representation




- order properties




# Advanced directories


an example:  
Venn diagram  
representation



## • query



## • navigate






# From advanced directories to advanced files

- in hierarchical filesystems, files have the same problems as directories: **rigid contents** (file contents vs directory contents)
- we have proposed a substitute to directory trees  
navigation and querying **among** files: **advanced dirs**
- the missing link:  
navigation and querying **inside** files: **advanced files**

# Hierachies, and their problems, are everywhere


- a file contents (symbolic)



- a latex file

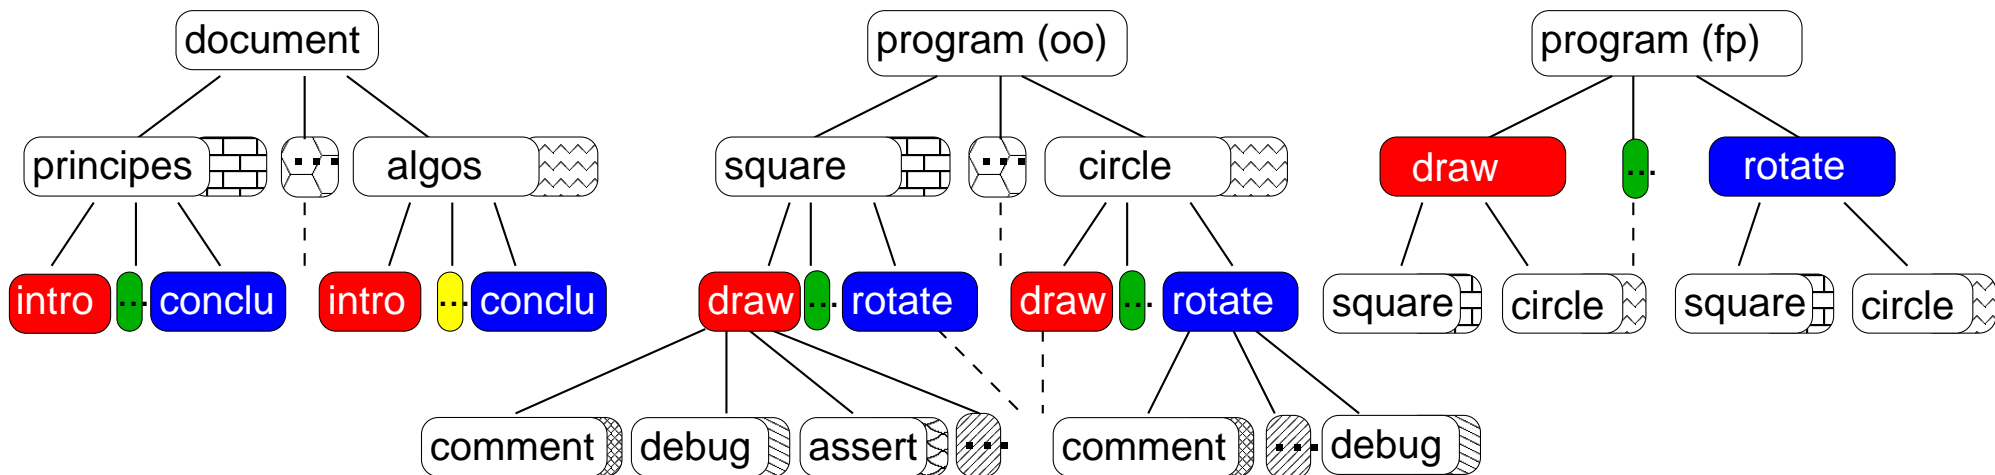


- file structure (symbolic)



- a c++ file

- a caml file



# Motivation

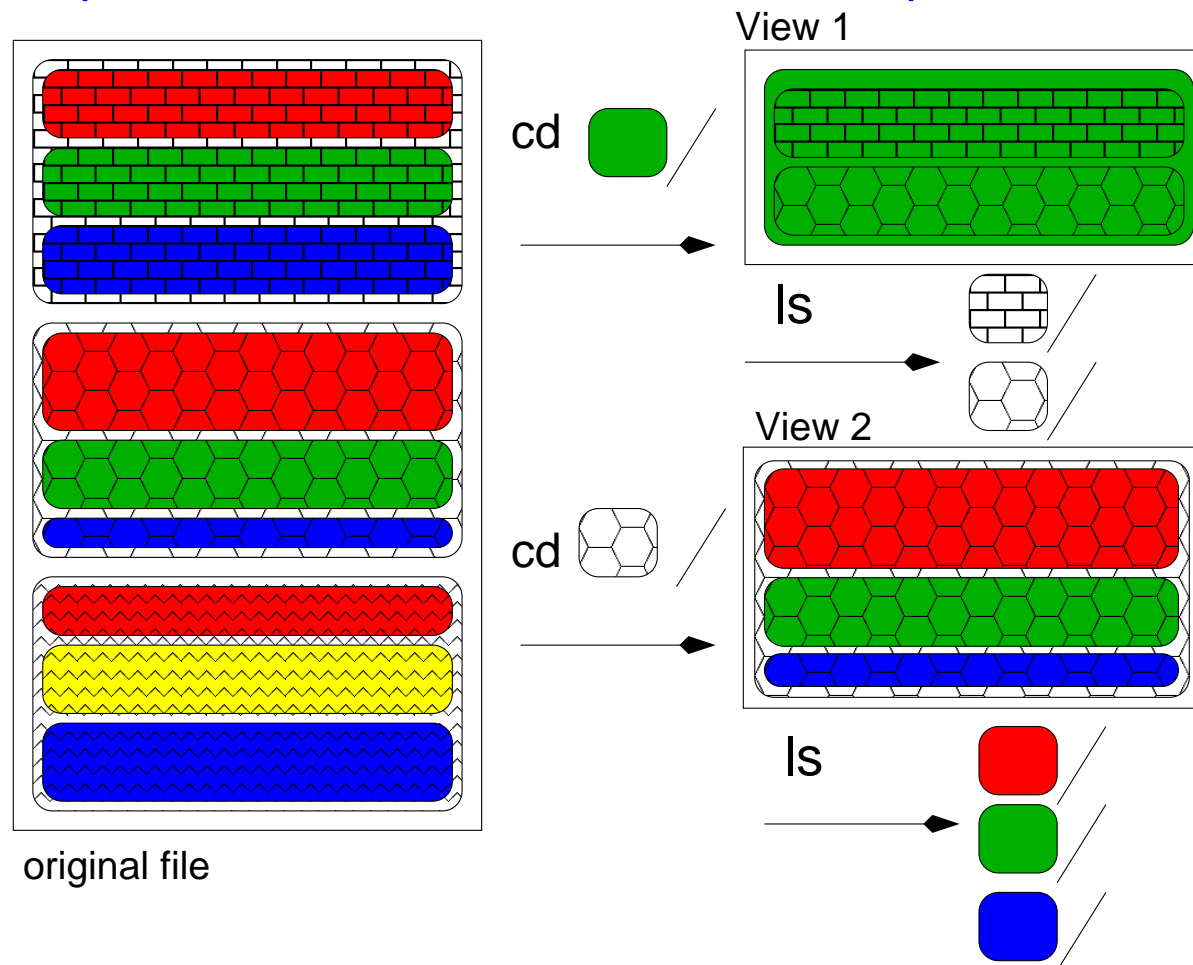
→ query,  
navigation,  
editing of points of view  
on file contents

better understanding  
easy focus on one task  
coherent change

advanced (virtual) files

# Advanced files

- assign multiple properties to parts of file
- query, navigate, view, update



# Demo

## Party Time!

# Specification

based on a previous work on **Logic Information Systems** (LIS)  
[Ferré&Ridoux, IP&M'2004]

(LFS is to LIS what O'Caml is to  $\lambda$ calcul)

# LFS content

- **a logic** —  $A \models B$ 
  - core deduction rules  $(a \wedge b \models a)$
  - axioms  $(music \models art)$
  - additional **logic engines**  $(size:4Mo45Ko \models size:>3Mo)$
- **information** — an attachment  $d$  (description) of a logic formula to every object  $o$  of a collection  $\mathcal{O}$  (**files** or **parts** of file)  
 $d(SanDiego.jpg) = port \wedge seaside \wedge USA \wedge size:4Mo45Ko$   
 $d(\mathbf{y} = \mathbf{x};) = var:x \wedge var:y \wedge function:f$ 
  - user gives properties  $(interesting, tested)$
  - LFS gives system properties  $(size:3Mo, owner:pad)$
  - additional **transducers** give properties  $(res:640x480, var:x)$

# Querying LFS

- paths are formulas
- **extension** — given a path  $pwd$ , the set of all objects that satisfy this path

$$ext(pwd) = \{o \in \mathcal{O} \mid d(o) \models pwd\}$$

$$ext(pwd) \approx ls -R \ pwd$$

*LosAngeles.jpg*  $\in ext(art)$  because

$$d(LosAngeles.jpg) = movie \wedge port \wedge USA \models movie \models art$$

- paths denote **directories** that denote **extensions**

(working directory  $\equiv$  working query)



# Navigating LFS

- to be a subdirectory of a directory — given a directory (denoting)  $O$ , every directory (denoting)  $O'$  such that  $O' \subset O$

$$Dirs(pwd) = \max_{\models} \{p \in \mathcal{P} \mid \emptyset \subset \text{ext}(pwd \wedge p) \subset \text{ext}(pwd)\}$$

*p refines pwd*

only «largest» subdirectories are relevant to navigation

(most relevant hints)

- to be a file of a directory — given a path  $pwd$ , to be in  $\text{ext}(pwd)$ , and in the extension of no subdirectory

$$Files(pwd) = \text{ext}(pwd) - \bigcup_{p \in Dirs(pwd)} \text{ext}(p)$$

$$Files(p) \cup Dirs(p) \approx ls \ p$$

# Security

- security properties = logic properties (secu:alice=rw)
- security rules = logic rules (rw  $\models$  r, group  $\models$  user)  
→ security in LFS (almost) for free (à la ACL and more)
- examples:
  - mv foo.c secu:alice=rw/secu:bob=r/
  - mv foo.txt secu:not(ridoux)=rw/
  - cd owner:pad/secu:group(student)=r/(only the owner of the file can change the properties of the file)

# Implementation

# Basic principles

to implement the specification at a reasonable cost

- to call  $\models$  is costly

avoid it by representing on disk a directed acyclic graph (DAG) of the properties: **logic cache**

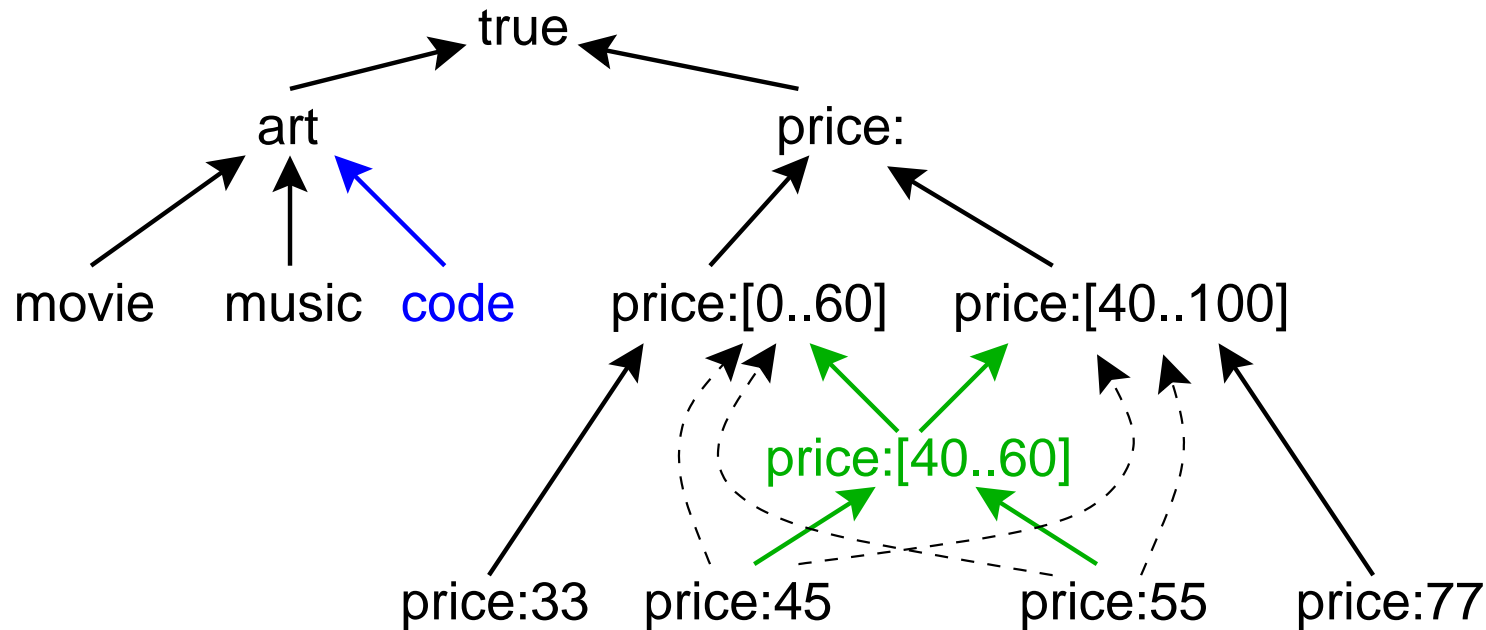
- to go through all objects ( $\{o \in \mathcal{O} \mid \dots\}$ ) and all properties ( $\{p \in \mathcal{P} \mid \dots\}$ ) is costly

avoid it by representing relation  $d$  as a table and **inverted** table on disk: **indexing**.

- to call part transducers is costly


define **synchronization points** for re-indexing

# Logic Cache

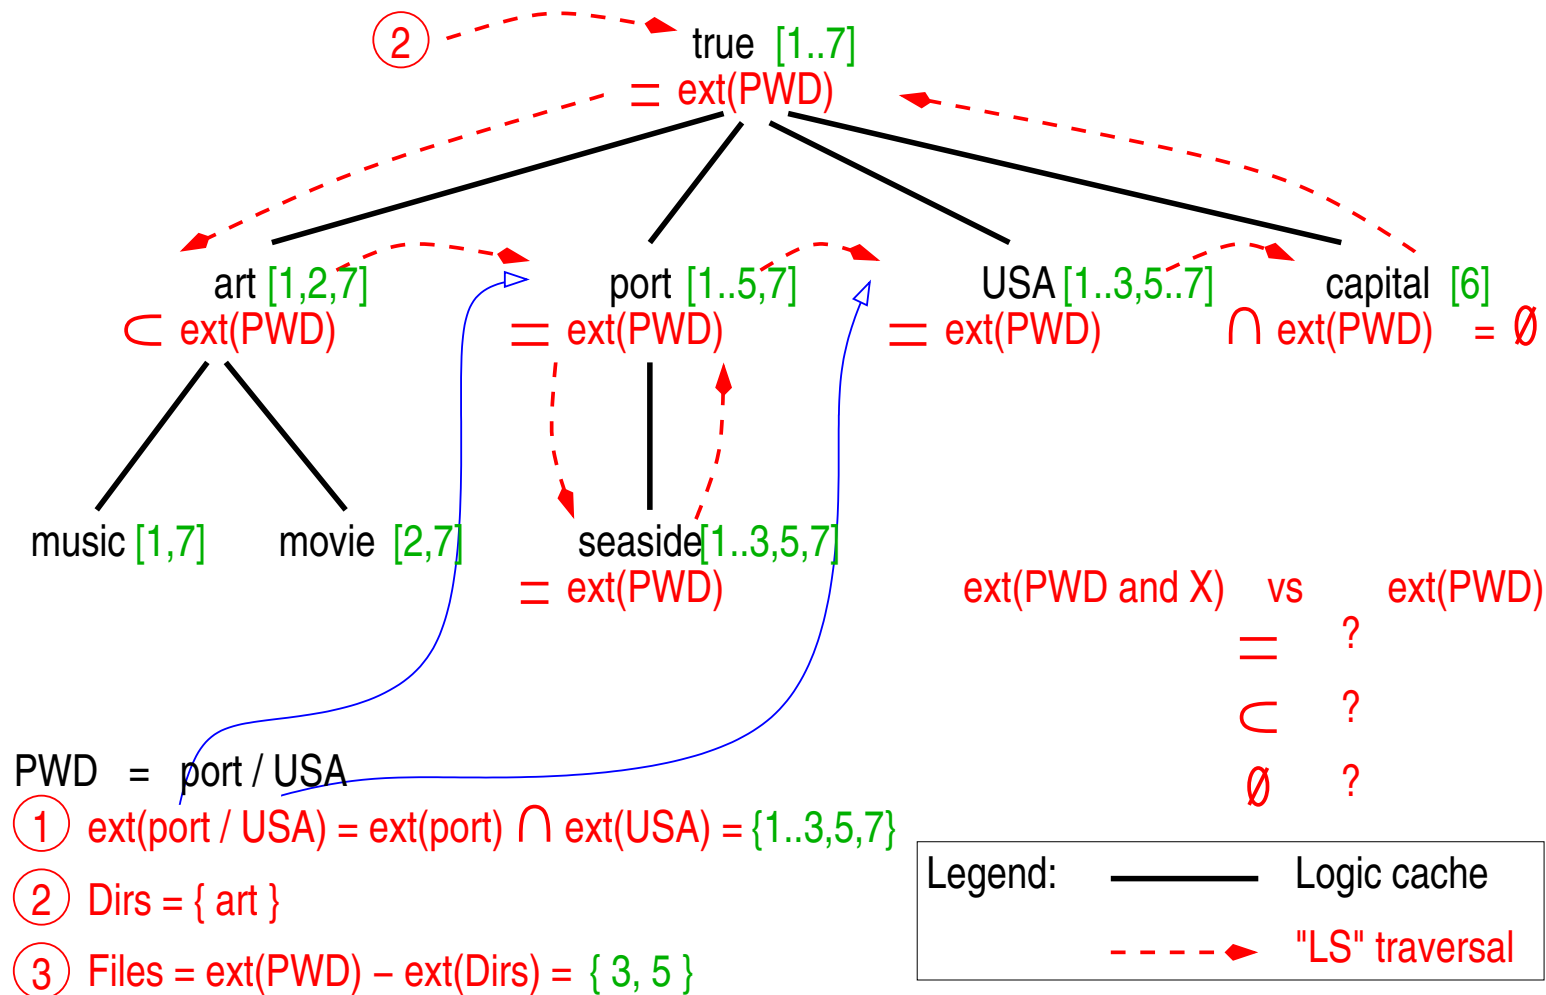


- `cd art/; mkdir code`
- `cd price:[40..60]/`

note that this is not the navigation structure  
(sub-property != sub-directory)

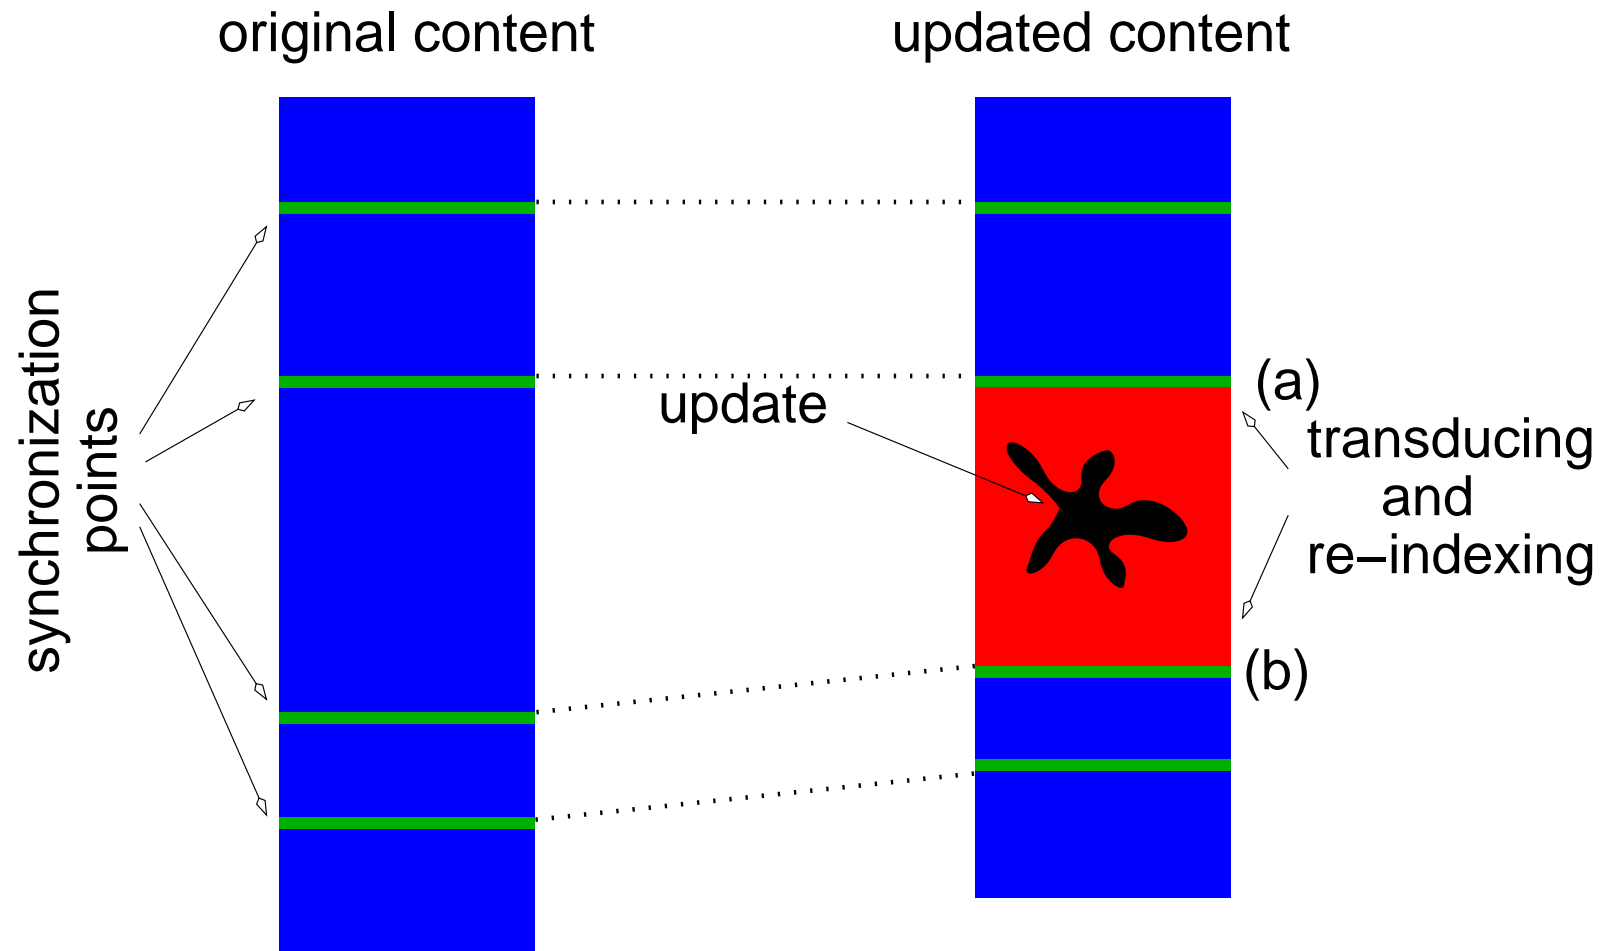


# Computing *Files* and *Dirs*



# Synchronization points

places in a file where a part transducer goes back to its **initial** state



# Evaluation


- **software** — a filesystem coded mostly in O'Caml
- **platform** — Linux kernel 2.4, with a 2Ghz Pentium 4, 750Mb RAM, and a 40 Gb IDE disk.



# Benchmark

- **data** — music, emails, man pages, latex, bibtex, ocaml sources
  - **context** —
    - up to 100 000 files ( $\approx 10$  prop per files)
    - up to 100 000 lines ( $\approx 10$  prop per parts)
  - **time** —
    - `ls`  $\approx 1$  sec (compute *Dirs* and *Files*)
    - `cd`  $\approx 1$  sec (may have to insert property in cache)
    - `cp`  $\approx 1$  sec (file transducer calls and indexing)
    - `cat`  $\approx 1$  sec (compute view content with *ext*)
    - `edit`  $\approx 1$  sec (part transducer calls and reindexing)
    - `cd parts` from 0.1 sec to 3 min (pure indexing)
- compatible with an interactive usage

# E-mails from linux kernel mailing list



# Related work

- advanced file systems

(SFS [Gifford et al.91], CATFS, BeFS [Giampaolo.93], Nebula [Bowman et al.94], HAC [Gopal & Manber.99])

—→ no navigation after arbitrary query

- information retrieval tools

(SuperBook [Remde et al.87], Scatter/Gather [Cutting et al.92], concept analysis [Lindig.95], Flamenco [Yee et al.03])

—→ navigation after query but poor query

- programming language IDEs

(Emacs [Stallman.81], Smalltalk [Goldberg et al.84], Omega [Linton.84], CIA [Chen et al.90], TuringTool [Cordy et al.90], Hyperspace [Tarr et al.99])

—→ limited query / navigation / view / updating inside files

# Contributions

- a **running** alternative to hierarchical file systems
- integration of **query** and **navigation**, with **manual** and **automatic** assignation and ordering of properties, for **files** and **parts** of files.  
(LFS is to LIS what O'Caml is to  $\lambda$ calcul)
- an architecture for **generic** services  
(many kind of files: **music**, **email**, **ocaml**, **perl**, **latex**, **bibtex**, ...)   
(many kind of logics: **boolean**, **intervals**, **regexp**, **type**, ...)
- a security model
- encouraging performances

→ **USENIX[2003]** (advanced dirs)

→ **USENIX[2005]** (advanced files)

# Further work

- more general kinds of parts: tokens, video segments, ...
  - [1] `cd scene:kung-fu|amour`
  - [2] `play movie.mpeg`
  - [3] `ls → Tom Cruise/ Monica Belluci/ JCVD/ ...`
- explore LFS usage
  - for software engineering (UML, AOP, slicing, versioning, ...)
  - for personal digital assistant (PDA)
- distribution
  - of computation (LFS algorithms are highly parallelizable)
  - of data (Semantic Web, peer-to-peer)
- relations/links between objects (`ln foo.o <compiled> foo.c`)
- improve performance (especially file creation)



# A LFS scenario (1/2)

---

mounting

```
[1] % mount /dev/hda1 /lfs/ ; cd /lfs/
```

---

taxonomy

```
[2] % mkdir art
```

```
[3] % cd art; mkdir music; mkdir movie; ... (adds music  $\models$  art,...)
```

---

context

```
[4] % cp /x/washington.jpg USA/capital/
```

```
[5] % cp /x/Boston.jpg seaside/USA/ (d(Boston.jpg) = seaside  $\wedge$  USA)
```

---

updating

```
[6] % cd USA/seaside/
```

```
[7] % mv Boston.jpg music/ (d(Boston.jpg) = music  $\wedge$  seaside  $\wedge$  USA)
```

---

navigating and querying

```
[8] % ls port/USA  $\longrightarrow$  art/ Miami.jpg SDiego.jpg
```

```
[9] % ls USA  $\longrightarrow$  art/ port/ capital/
```

```
[10] % ls !USA  $\longrightarrow$  Hamburg.jpg
```

```
[11] % cd port|art/!(USA&FR)/  $\longrightarrow$  ...
```

## A LFS scenario (2/2)

```
[1] % cp /x/foo.c
/lfs/project:foobar/
[2] % cp /x/trans_C.exe
/lfs/parts_transducer:c/
[3] % cd project:foobar/
[4] % cat foo.c
```

```
int f(int x) {
    int y;
    assert(x > 1);
    y = x;
    printf("x = %d", x);
    return y * 2
}
int f2(int z) {
    return z * 4
}
```

```
[5] % cd parts; ls
function:f/ function:f2/
var:x/ var:y/ var:z/
debug/ assert/
foo.c
[6] % cd function:f/-
!(debug|assert)/
[7] % cat foo.c
int f(int x) {
    int y;
    .....:1
    y = x;
    .....:2
    return y * 2
}
.....:3
```



# LFS organization

*movie*  $\models$  *art*    *music*  $\models$  *art*    ...

and

<i>Dirs</i> ←	art	music	movie	port	seaside	USA	capital
Boston.jpeg	x	x		x	x	x	
LosAngeles.jpeg	x		x	x	x	x	
Miami.jpeg				x	x	x	
Hamburg.jpeg				x			
SanDiego.jpeg				x	x	x	
Washington.jpeg						x	x
NewYork.jpeg	x	x		x	x	x	


*Files* →

## indexing

	function:f	function:f2	var:x	var:y	var:z	debugging	specification
1	X		X				
2	X			X			
3	X		X				X
4	X		X	X			
5	X		X			X	
6	X			X			
7	X						
8		X			X		
9		X			X		
10		X					

# updating LFS

- **re-indexing** — updating a view  $\rightarrow$  updating the object $\times$ property matrix,
- **marks of missing parts** — . . . . . : X

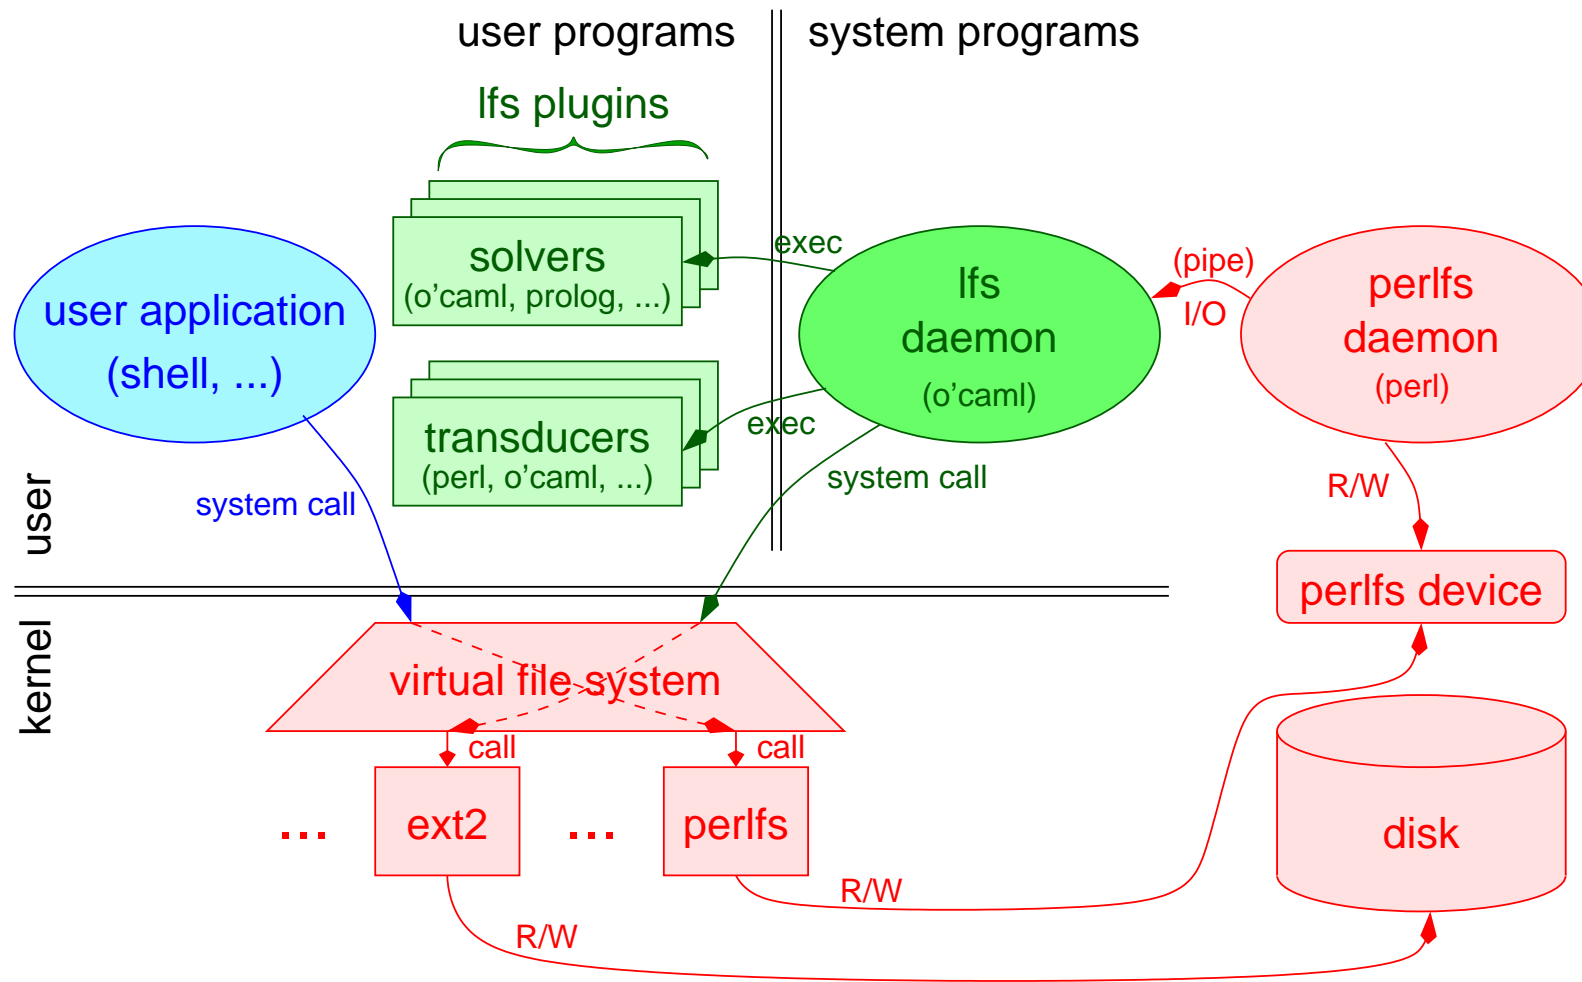


to know where to insert new parts

# Semantics of LFS operations

- **readdir(path)** — lists  $Files(path) \cup Dirs(path)$  (ls path)
- **lookup(name,path)** — checks  $name \in Files(path) \cup Dirs(path)$
- **create(name,path)** — adds  $d(name) = path$   
(touch path/name)
- **mkdir(name,path)** — adds  $name \models path$  (mkdir path/name)
- **file operations** — as usual (open, read, write, ...)
- ...

# LFS software architecture



# Do optimisations optimise ?

